Vaibhav Hiwase

AI/ML Engineer & Technical Lead — LLM Applications, Evaluation & Production AI Platforms

Nagpur, India · (+91) 8999382164 · hiwase.vaibhav@gmail.com · linkedin.com/in/vaibhav-hiwase

19 October 2026

[Hiring Manager Name]
[Company Name]
[City, Country]

Dear [Hiring Manager],

I am applying for the [Role Title] role at [Company]. I own two production AI systems on Palantir Foundry and AIP, and I bring a specific discipline: taking LLM systems past the demo stage and keeping them dependable once people rely on them every day. I measure every change before and after it ships.

At Aptia Group, I build a GenAI assistant that answers employees' benefits questions. I raised the accuracy of its coverage answers from 68.9% to 97.9% by replacing about 30 rigid lookup tables with a principle-based prompt on Gemini Flash. The change also cut token cost by 28% and worst-case response time by 48%. I rebuilt the classifier that sorts each question into one of 26 topics, which now routes 99.4% of test questions correctly, and reduced made-up answers by having the assistant cite its sources and run an 11-point self-check.

Evaluation is the part of this work I care about most. I designed a 532-question test suite from real plan documents, with an AI judge that must agree with human experts at least 90% of the time. Business users wrote 388 of those questions, because the people who define what "correct" means should own that standard. A separate check found 85 of 182 business test questions that the source documents could not answer.

Beyond the assistant, I automate document verification and run the platform. I added a readiness check to 22 automations so requests no longer get stuck when a person's data is not yet available. I also built load tests that measure response times at up to 250 times baseline load, and a deployment-checking tool, backed by 158 unit tests, that compares resources across five environments. As Platform Admin, I provide L3 production support.

Before Foundry, I led contract-document AI at TechnoMile. I raised extraction accuracy from 58% to 90% across seven contract domains and split one large service into four smaller services that ran 40% faster.

I would welcome the chance to discuss [specific team, product or problem at Company] and where my experience would help most. Thank you for your time and consideration.

Sincerely,

Vaibhav Hiwase